

# INTRODUCTION TO pandas

**pandas** objects can be thought of as “enhanced” versions of **numpy** arrays in which the rows and columns are identified with labels rather than integers. Basics **pandas** objects:

- **Series**: 1D array of data with an index object (labels).

```
In [4]: import pandas as pd

column = pd.Series([0.25, 0.5, 0.75, 1.0],
                    index=['one', 'two', 'three', 'four'])
column

Out[4]: one      0.25
        two      0.50
        three    0.75
        four     1.00
        dtype: float64

In [6]: column['four']

Out[6]: 1.0
```

Each series has a “values” component and an “index” component. Series come with the same built-in functions, like mean, min, std, as **numpy** arrays.

# INTRODUCTION TO pandas

**pandas** objects can be thought of as “enhanced” versions of **numpy** arrays in which the rows and columns are identified with labels rather than integers. Basics **pandas** objects:

■ **Series:** 1D array of data with an index object (labels).

```
In [7]: column.index
```

```
Out[7]: Index([u'one', u'two', u'three', u'four'], dtype='object')
```

```
In [11]: column.values
```

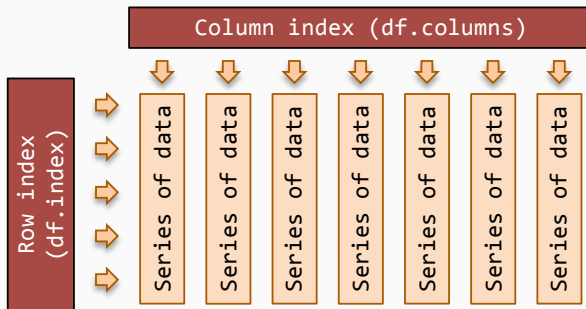
```
Out[11]: array([ 0.25,  0.5 ,  0.75,  1.  ])
```

Each series has a “values” component and an “index” component. Series come with the same built-in functions, like mean, min, std, as **numpy** arrays.

# INTRODUCTION TO pandas

**pandas** objects can be thought of as “enhanced” versions of **numpy** arrays in which the rows and columns are identified with labels rather than integers. Basics **pandas** objects:

- **DataFrame**: 2D table of data with column and row index objects (labels).



Each column in the data frame is a *series*.

# GETTING STARTED WITH pandas

We can create a data frame from columns (series objects):

```
In [15]: column_1 = pd.Series(range(4),  
                               index=['one', 'two', 'three', 'four'])  
  
column_2 = pd.Series(range(4, 8),  
                      index=['one', 'two', 'three', 'four'])  
  
table = pd.DataFrame({'col_1': column_1,  
                      'col_2': column_2})  
  
table
```

Out[15]:

	col_1	col_2
one	0	4
two	1	5
three	2	6
four	3	7

## GETTING STARTED WITH pandas

We can import tabular data in a csv file into a data frame:

```
In [16]: df = pd.read_csv('dataset_HW0.txt')  
df
```

Out[16]:

	birth_weight	femur_length	mother_age
0	2.969489	1.979156	16
1	4.038963	3.555681	16
2	5.302643	3.385633	15
3	6.086107	4.495427	17

We should start by getting a rough sense of what's in the data

## ■ The indices of your data frame:

```
In [32]: df.columns
Out[32]: Index([u'birth_weight', u'femur_length', u'mother_age'], dtype='object')

In [5]: df.columns.values
Out[5]: array(['birth_weight', 'femur_length', 'mother_age'], dtype=object)

In [6]: df.index
Out[6]: Int64Index([ 0, 1, 2, 3, 4, 5, 6, 7, 8, 9,
...,
390, 391, 392, 393, 394, 395, 396, 397, 398, 399],
dtype='int64', length=400)

In [7]: df.index.values
Out[7]: array([ 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12,
13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25,
26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38,
39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51,
```

The `.index` and `.columns` attributes give access to the index objects of rows and columns (resp).

We should start by getting a rough sense of what's in the data

### ■ The shape of your data frame:

```
In [23]: df.shape
```

```
Out[23]: (400, 3)
```

```
In [10]: len(df.index)
```

```
Out[10]: 400
```

## GETTING TO KNOW YOUR DATAFRAME

We should start by getting a rough sense of what's in the data

### ■ The first entries in your data frame:

```
In [25]: df.head(n=5)
```

```
Out[25]:
```

	birth_weight	femur_length	mother_age
0	2.969489	1.979156	16
1	4.038963	3.555681	16
2	5.302643	3.385633	15
3	6.086107	4.495427	17
4	5.749260	4.017437	16

The `.head()` function returns a (row-wise) truncated version of your data frame!



## GETTING TO KNOW YOUR DATAFRAME

We should start by getting a rough sense of what's in the data

### ■ A summary of your data frame:

```
In [29]: df.describe()
```

```
Out[29]:
```

	birth_weight	femur_length	mother_age
count	400.000000	400.000000	400.000000
mean	6.104070	3.827591	27.060000
std	1.097011	0.853577	10.349840
min	2.967426	0.479154	15.000000
25%	5.429120	3.281786	17.750000
50%	6.110025	3.817888	25.000000
75%	6.839935	4.351204	34.250000
max	9.021942	6.648730	49.000000

The `.describe()` function returns all the descriptive stats for each column as a data frame object!

### ■ Accessing a column by label:

```
In [34]: type(df['birth_weight'])
```

```
Out[34]: pandas.core.series.Series
```

```
In [35]: df['birth_weight']
```

```
Out[35]: 0      2.969489  
         1      4.038963  
         2      5.302643  
         3      6.086107  
         4      5.749260  
         5      6.049903
```

You can access a column by it's column name or position (you can also access a *list* of columns)!

## ■ Accessing the values of column:

```
In [36]: df['birth_weight'].values
```

```
Out[36]: array([ 2.9694893 ,  4.03896294,  5.30264328,  6.08610661,  5.74926036,  
                 6.04990317,  5.42681579,  6.23910323,  5.34504952,  4.16297458,  
                 5.27487188,  5.57627684,  5.49364519,  6.66031745,  4.79466787,  
                 5.98546786,  4.62521954,  5.60683336,  4.52477222,  6.3162985 ,  
                 5.5922901 ,  6.23730155,  5.19645533,  4.61051962,  4.38347209,  
                 5.00708476,  4.10801732,  5.18226899,  3.91916625,  5.8955964 ,
```

You can access a column by it's column name or position (you can also access a *list* of columns)!

### ■ Accessing columns by position:

```
In [63]: df[[0, 1]]
```

```
Out[63]:
```

	birth_weight	femur_length
0	2.969489	1.979156
1	4.038963	3.555681
2	5.302643	3.385633
3	6.086107	4.495427
4	5.749260	4.017437
5	6.049903	4.378892
6	5.426816	2.851801

You can access a column by it's column name or position (you can also access a *list* of columns)!

### ■ Accessing a row by position:

```
In [46]: type(df.iloc[0])
```

```
Out[46]: pandas.core.series.Series
```

```
In [47]: df.iloc[0]
```

```
Out[47]: birth_weight      2.969489  
         femur_length      1.979156  
         mother_age       16.000000  
         Name: 0, dtype: float64
```

You can access a column by it's row name or position!

## ■ Accessing a row by label:

```
In [52]: table
```

```
Out[52]:
```

	col_1	col_2
one	0	4
two	1	5
three	2	6
four	3	7

```
In [54]: type(table.loc['one'])
```

```
Out[54]: pandas.core.series.Series
```

```
In [53]: table.loc['one']
```

```
Out[53]: col_1    0
         col_2    4
         Name: one, dtype: int64
```

You can access a column by it's row name or position!

Filtering works very much like with `numpy` arrays!

```
In [57]: df[(df['mother_age'] > 18) & (df['mother_age'] < 35)]
```

```
Out[57]:
```

	birth_weight	femur_length	mother_age
100	6.904530	4.164637	34
101	8.096642	4.536759	22
102	8.165373	5.507030	20
104	6.255286	3.769024	19
105	6.515220	5.568954	23
106	6.464462	3.310628	25
107	6.579616	3.670224	20
108	7.171024	5.159946	24